

Unit 5

Building blocks of CNNs:

1. Input Layer

- **Purpose:** Accepts the image as input.
- **Format:** 3D matrix \rightarrow Height \times Width \times Channels
(e.g., $224 \times 224 \times 3$ for RGB image)

✓ 2. Convolutional Layer

- **Purpose:** Extracts features like edges, textures, patterns.
- **Operation:** Applies filters (kernels) that slide over the image.
- **Output:** Feature Maps.

✓ 3. Activation Function (ReLU)

- **Purpose:** Adds non-linearity to the model.
- **Function:** $\text{ReLU}(x) = \max(0, x)$
- **Effect:** Helps model learn complex patterns.

✓ 4. Pooling Layer

- **Purpose:** Reduces spatial dimensions (height, width).
- **Types:**
 - **Max Pooling:** Takes maximum value.
 - **Average Pooling:** Takes average value.
- **Benefits:**
 - Reduces computation.
 - Controls overfitting.
 - Adds translation invariance.

✓ 5. Padding

- **Purpose:** Preserves image size after convolution.
- **Types:**
 - **Valid Padding:** No padding, output shrinks.
 - **Same Padding:** Adds padding to keep size same.

✓ 6. Stride

- **Definition:** Step size of the filter as it moves.
- **Effect:**
 - **Stride = 1:** Normal scan.
 - **Stride > 1:** More aggressive reduction in size.

✓ 7. Fully Connected (Dense) Layer

- **Purpose:** Combines features and performs classification.
- **Structure:** Each neuron connects to all neurons in previous layer.

✓ 8. Output Layer (Softmax)

- **Purpose:** Converts outputs into probabilities.
- **Function:** Used for multi-class classification.
- **Example:** 10 units for classifying digits 0–9.


✓ 9. Other Components

- **Batch Normalization:** Normalizes layer inputs, improves training.
- **Dropout:** Randomly disables neurons during training to prevent overfitting.
- **Flatten Layer:** Converts 2D/3D feature maps into 1D vector before fully connected layer.

Input Image \rightarrow Convolution \rightarrow ReLU \rightarrow Pooling \rightarrow Convolution \rightarrow ReLU \rightarrow Pooling \rightarrow Flatten \rightarrow Fully Connected \rightarrow Softmax Output

Architecture of a CNN

A CNN is made up of multiple layers that work together to extract and learn features from images. Here's a step-by-step breakdown of a **typical CNN architecture**:



1. Input Layer

- **Role:** The input layer takes in the raw pixel values of an image.
- **Structure:** The image is represented as a **3D matrix**:
 - Height \times Width \times Channels
 - Examples:
 - Grayscale image: $28 \times 28 \times 1$
 - RGB image: $224 \times 224 \times 3$
- This layer does **not learn** any parameters; it just passes the data to the next layer.

2. Convolutional Layers

- **Role:** Extract local patterns or features from the input image.
- **How It Works:**
 - Applies a set of **learnable filters (kernels)**.

- Each filter slides (convolves) across the image.
- Computes **dot products** between the filter and regions of the image.
- Produces a **feature map** that indicates where a particular feature appears.
- **Filters** detect:
 - Low-level features: edges, corners
 - High-level features: textures, objects (in deeper layers)
- **Parameters learned:** weights of filters.

↪ 3. Activation Function (ReLU)

- **Role:** Introduce **non-linearity** into the network.
- **Why Needed:** Without non-linear activations, a CNN would be equivalent to a linear model.
- **ReLU (Rectified Linear Unit):**

$$\text{ReLU}(x) = \max(0, x) \quad \text{ReLU}(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

- Converts all negative values to 0.
- Keeps positive values unchanged.
- **Benefits:**
 - Faster computation
 - Avoids vanishing gradient issues (better than sigmoid/tanh)

▼ 4. Pooling Layer

- **Role:** Reduce spatial dimensions (height and width) of feature maps.
- **Common Types:**
 - **Max Pooling:** Takes the maximum value in a patch.
 - **Average Pooling:** Takes the average value.
- **Why It's Important:**
 - Reduces number of parameters and computations.
 - Makes features more robust to translation/rotation.
 - Prevents overfitting.
- **Typical Configurations:** Pool size = (2×2), Stride = 2

□ 5. Padding

- **Role:** Add extra border (usually zeros) around the input matrix.
- **Why Use Padding?**
 - Prevents shrinking of output after each convolution.
 - Allows **filters to reach the edges** of the image.
- **Types:**
 - **Valid Padding:** No padding → output size reduces.
 - **Same Padding:** Padding added → output size stays the same.

↔ □ 6. Stride

- **Definition:** Number of pixels the filter moves at each step.
- **Effects:**
 - **Stride = 1:** Filters move one pixel at a time → high detail.
 - **Stride > 1:** Filters move faster → reduced output size.
- **Impact on Architecture:**
 - Larger strides = smaller outputs = less computation.

□ 7. Flatten Layer

- **Role:** Converts multidimensional feature maps into a **1D vector**.
- This vector is then used as input to the Fully Connected Layer.
- **Why Needed?:**
 - Dense layers expect 1D input.
 - Bridges the transition from convolutional (spatial) to classification (vector) tasks.

∞ 8. Fully Connected (Dense) Layer

- **Role:** Learn **high-level features** and perform reasoning based on extracted features.
- **Structure:** Every neuron is connected to all neurons in the previous layer.
- **Output:** Feature vector for final decision-making.
- **Example:** In digit classification, 10 neurons for digits 0–9.

□ 9. Output Layer

- **Role:** Generate the final prediction.
- **Common Activation Functions:**
 - **Softmax:** Multi-class classification (outputs probability for each class).

$$\text{Softmax}(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

- **Sigmoid:** Binary classification (outputs value between 0 and 1).

$$\text{Sigmoid}(x) = \frac{1}{1 + e^{-x}}$$

1. Convolution Layer

Definition:

A **convolutional layer** is the fundamental building block of a CNN, designed to automatically and adaptively learn **spatial hierarchies of features** from input images.

How it works:

- It uses a **filter (kernel)** — a small matrix (e.g., 3×3 or 5×5) — which slides over the input image.
- At each position, it performs **element-wise multiplication** between the filter and the corresponding input patch, then sums the results to form a single number.
- This operation results in a new matrix called a **feature map** or **activation map**.

Purpose:

- Detects low-level features (edges, corners) in initial layers and high-level features (shapes, textures) in deeper layers.
- Reduces the number of parameters compared to fully connected layers.

Formula:

$$O(i, j) = \sum_m \sum_n I(i + m, j + n) \cdot F(m, n)$$

2. Pooling Layers

Pooling layers are used to **reduce the spatial dimensions** (height and width) of the feature maps, while **retaining important features**.

◆ A. Max Pooling

Definition:

Selects the **maximum value** from each patch (e.g., 2×2) of the feature map.

Purpose:

- Preserves the most prominent feature in a region.
- Adds **translation invariance** (small shifts in input won't affect output).
- Helps reduce overfitting and computational cost.

◆ B. Average Pooling

Definition:

Computes the **average value** from each patch.

Purpose:

- Provides a smoother version of the feature map.
- Less aggressive than max pooling.
- Can be useful when presence of a feature is less important than its average strength.

◆ C. Min Pooling

Definition:

Selects the **minimum value** from each patch.

Purpose:

- Captures the least-activated regions.
- Useful in specific tasks like anomaly or defect detection where the absence or weakness of a feature is meaningful.

3. Padding

Definition:

Padding refers to the **addition of extra pixels** (usually zeros) around the border of the input image or feature map before applying convolution.

Types:

- **Same Padding:** Pads input so that output has **same spatial dimensions**.
- **Valid Padding:** **No padding**, resulting in reduced output size.

Purpose:

- Allows the network to preserve edge information.
- Ensures deeper networks do not excessively shrink spatial dimensions.
- Controls the size of the output feature maps.

Padding Formula (Same):

$$\text{For a filter size } k, \text{ padding } p = \left\lfloor \frac{k-1}{2} \right\rfloor$$

◆ 4. Strided Convolutions

Definition:

Stride refers to the **number of pixels by which the filter moves** across the input image.

How it affects output:

- **Stride = 1:** Filter moves one pixel at a time — results in detailed feature maps.

- **Stride > 1:** Filter skips pixels — results in down sampling and faster computation.

Purpose:

- To control the spatial resolution of the output.
- Acts as a **learned down sampling technique**.
- Reduces the number of computations and parameters.

Output Size Formula: For input size N , filter size F , padding P , and stride S :

$$\text{Output size} = \left\lfloor \frac{N - F + 2P}{S} \right\rfloor + 1$$

Convolutions over Volumes in CNNs

In the context of **Convolutional Neural Networks (CNNs)**, **convolutions over volumes** refer to the process of applying convolutional operations not just to 2D images (height \times width) but to 3D inputs — known as **volumes**. This 3D input can be represented as a matrix of **height \times width \times depth**, where **depth** typically refers to the number of **channels** in the image (for example, 3 channels for RGB images).

1. What is a Volume?

A **volume** in CNNs is a 3D matrix (height \times width \times depth), where:

- **Height** refers to the vertical dimension of the input image.
- **Width** refers to the horizontal dimension of the input image.
- **Depth (Channels)** refers to the number of feature maps or channels. For example:
 - A **grayscale image** has only one channel (depth = 1).
 - A **RGB image** has three channels (depth = 3).
 - After applying convolution in earlier layers, depth might increase due to multiple feature maps.

2. Applying Convolution Over Volumes

When we apply convolutions over volumes, each **filter** (or **kernel**) is designed to slide across the entire 3D volume, instead of just 2D images. This means that the filter also has a **depth dimension** (in addition to height and width), and it operates over a **3D region** in the input volume.

How It Works:

1. Filter/Kernel Dimensions:

For a volume, a filter will have a size such as (**f** \times **f** \times **d**), where:

- **f \times f** is the spatial dimension (height \times width) of the filter.
- **d** is the depth of the filter, which must match the depth of the input volume (for an RGB image, this is 3).

2. Sliding the Filter:

The filter moves across the input volume both **horizontally** and **vertically**, considering the **depth** as well. At each position, a **3D dot product** is performed between the filter and the corresponding region of the input volume.

3. Output Volume:

The result is a **3D output volume** (feature map), where each depth slice represents a different feature detected by the filter. The depth of the output corresponds to the number of filters applied.

Mathematical Representation:

- For an input volume V of size (height \times width \times depth) and a filter W of size ($f \times f \times \text{depth}$), the operation at a given location (i, j) in the output volume is:

$$\text{Output}(i, j) = \sum_{m=1}^f \sum_{n=1}^f \sum_{k=1}^{\text{depth}} V(i+m, j+n, k) \cdot W(m, n, k)$$

- This operation will be repeated as the filter slides across the height and width of the input volume.

Output from Convolutions Over Volumes

After applying a set of filters, the output will be a **new feature map (3D volume)**. The number of channels in the output volume corresponds to the number of filters used.

- **Output Size:** The size of the output feature map depends on:
 - The **input size**.
 - The **filter size**.
 - The **stride**.
 - The **padding**.

Formula for the output volume size:

$$\text{Output height} = \left\lfloor \frac{(H - F + 2P)}{S} \right\rfloor + 1$$

Where:

- H = Height of input volume
- F = Filter size
- P = Padding
- S = Stride

Softmax Regression

Softmax regression, also known as **multinomial logistic regression**, is a generalization of logistic regression to multi-class classification problems. It is used to model the probability distribution of a list of potential outcomes.

Working Mechanism:

- **Normalization:** Softmax regression transforms the raw output scores (logits) from the neural network into probabilities by normalizing them. This is done through the **Softmax function**, which converts a vector of real numbers into a probability distribution.
- **Mathematical Formulation:**

Mathematical Formulation:

Given a vector of scores z_1, z_2, \dots, z_K (where K is the number of classes), the Softmax function is applied as:

$$\text{Softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

This formula ensures that the output for each class z_i lies in the range $[0,1]$ and that the sum of all outputs equals 1.

Key Features:

- **Multi-Class Classification:** Unlike binary logistic regression, which is used for two-class problems, **Softmax regression** can handle **multiple classes** (e.g., classifying an image as one of many categories like dog, cat, or bird).
- **Output:** The output of the Softmax function is a **probability distribution**, meaning that the predicted values will be between 0 and 1, and the sum of the probabilities across all classes will equal 1.

Advantages:

- **Multi-Class Capability:** Softmax regression avoids binary classification limitations and allows for classification problems with more than two classes.
- **Probabilistic Interpretation:** The output can be interpreted as the probability of the input belonging to each class, which is useful for applications like **multi-class classification** and **uncertainty estimation**.

Example:

For a neural network with 3 classes, Softmax regression might convert raw output scores (logits) like:

$$\text{Logits} = [2.0, 1.0, 0.1]$$

into probabilities (after applying Softmax):

$$\text{Softmax Output} = [0.659, 0.242, 0.099]$$

This means the model predicts the first class with 65.9% probability, the second class with 24.2%, and the third class with 9.9%.

Deep Learning Frameworks

Deep learning frameworks are powerful tools that simplify the process of building, training, and deploying deep learning models. These frameworks provide high-level abstractions, optimized performance, and easy-to-use tools, which help researchers and developers implement complex neural networks. Some of the most popular deep learning frameworks include **TensorFlow**, **PyTorch**, **Keras**, and **Theano**.

1. TensorFlow

TensorFlow is an open-source deep learning framework developed by Google in 2015 for numerical computation and machine learning applications. It is one of the most popular frameworks, especially for production-level deep learning projects.

Working:

TensorFlow uses a **dataflow graph** to define the structure of computation. It represents the computation in the form of a graph, where each node is a mathematical operation, and edges represent data (tensors). TensorFlow allows users to prepare a flowchart, where inputs (as tensors) are processed through these operations to generate outputs.

Applications:

- **Text-based Applications:** TensorFlow is widely used in natural language processing (NLP) tasks like **language detection** and **sentiment analysis**.

- **Image Recognition:** TensorFlow powers various image-related applications such as **motion detection**, **facial recognition**, and **photo clustering**.
- **Video Detection:** It is also used in **real-time object detection** from video feeds, enabling tracking and motion detection.

2. PyTorch

PyTorch is another popular deep learning framework, developed by Facebook's AI Research Lab (FAIR). It is known for its dynamic computational graph and strong support for GPU acceleration.

Working:

PyTorch uses a **dynamic computational graph**, which is built at runtime. This means you can change the graph structure on the fly, making it easier to debug and modify models. PyTorch also integrates with Python's core programming concepts, including loops and conditional statements, which simplifies model building and training.

Applications:

- **Weather Forecasting:** PyTorch is used to predict weather patterns based on historical data and real-time inputs.
- **Text Auto-detection:** It powers **auto-suggestions** in search engines (like Google) and other **text-based prediction** tasks.
- **Fraud Detection:** PyTorch helps in detecting anomalies and frauds, such as **credit card fraud** by recognizing unusual patterns in transaction data.

3. Keras

Keras is a high-level neural network API, written in Python, designed to enable easy and fast prototyping. Initially developed as a standalone library, Keras now runs on top of other deep learning frameworks like TensorFlow and Theano.

Working:

Keras provides an easy-to-use API for defining neural networks. It acts as a wrapper for low-level libraries (e.g., TensorFlow or Theano) to simplify model construction. Its modular approach allows for quick experimentation and testing of deep learning models, making it ideal for researchers and engineers who want to prototype quickly.

Applications:

- **Smartphone Applications:** Keras is used to develop machine learning models that power deep learning applications on smartphones.
- **Healthcare:** It is applied in the **prediction of heart diseases**, where models can alert healthcare providers about potential risks.
- **Face Mask Detection:** During the COVID-19 pandemic, Keras was used to develop models

for detecting whether a person is wearing a mask in real-time.

4. Theano

Theano is one of the earliest deep learning frameworks and is widely known for its ability to optimize mathematical computations. Developed by the **MILA (Montreal Institute for Learning Algorithms)**, Theano is no longer actively developed, but its legacy continues to influence modern frameworks like Keras.

Working:

Theano works by defining **mathematical expressions** and performing **symbolic differentiation** for training deep learning models. It can be used to run computations on both CPUs and GPUs, optimizing for the hardware used. Theano enables automatic differentiation, a key feature for training neural networks.

Applications:

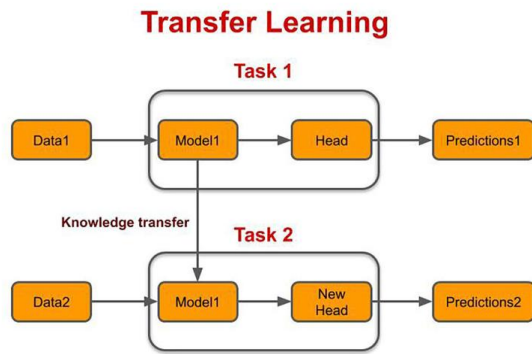
- **Neural Network Implementation:** Theano was widely used in implementing neural networks, particularly in research.
- **Optimization in Large-Scale Applications:** Companies like **IBM** have used Theano for developing large-scale neural networks and improving performance.
- **Scientific Computing:** Theano's optimization techniques are useful for working with large datasets and performing complex mathematical computations in deep learning.
- model and fine-tune it for your task.

Training and testing on different distributions, Bias and Variance with mismatched data distributions,

What Is Transfer Learning?

Transfer learning is a technique in machine learning where a model developed for a specific task is reused or adapted to solve a related but different task. Instead of training a model from scratch, transfer learning leverages the knowledge gained by a pre-trained model on one task to improve generalization on a new task. Essentially, the model transfers learned features, weights, or patterns to a different, often smaller or more specialized, dataset or

task.



For example, if a model is trained to classify images of animals, the knowledge it gained about shapes, textures, and edges can be transferred to a new model that classifies medical images, even if the new dataset is smaller and lacks a large amount of labeled data.

How Transfer Learning Works

1. **Pre-trained Model:**
 - A model that has already been trained on a large dataset for a particular task is used as the starting point.
 - This model has learned a rich set of features (e.g., edges, shapes, textures for image tasks) that are generally applicable to many different tasks.
2. **Fine-tuning:**
 - Fine-tuning involves slightly adjusting the weights of the pre-trained model on the new dataset. This helps the model specialize in the new task without starting from scratch.
 - In some cases, the early layers (which learn generic features) are frozen (not updated during training), and only the later layers are trained for the new task.
3. **Feature Extraction:**
 - Instead of fine-tuning the entire model, the pre-trained model can also be used as a **feature extractor**. The features learned in the earlier layers are used as input to a new classifier or model tailored to the new task.

Why Use Transfer Learning?

1. **Saves Training Time:** Pre-trained models reduce training time significantly by already having learned general features, speeding up the process for new tasks.
2. **Reduces Need for Large Datasets:** Transfer learning allows the use of smaller datasets, as the model has already learned important features from the original task, minimizing the need for vast labeled data.
3. **Improves Performance:** The pre-trained model brings valuable knowledge, often improving performance on the new task, especially when data is limited.
4. **Applicability in Specialized Fields:** Transfer learning is useful in domains like NLP and Computer Vision, where labeled data is scarce,

allowing general models to be adapted for specific tasks such as sentiment analysis or medical image analysis.


• Example 1: Image Classification:

- You want to classify photos of dogs and cats, but you have a very small dataset (e.g., 100 images). Instead of training a neural network from scratch, you start with a model pre-trained on ImageNet (which has millions of images across thousands of classes) and fine-tune it on your small dataset.

When to Use Transfer Learning

1. **Lack of Training Data:** If there is insufficient labeled data for the task, transfer learning allows you to use a pre-trained model and fine-tune it with a smaller dataset.
 2. **Existing Pre-trained Network:** If a pre-trained model exists for a similar task (e.g., image classification), you can reuse the learned features to speed up training and improve performance.
 3. **Same Input Type:** Transfer learning works well when both tasks use the same type of input, such as images or text. For example, using an image classifier trained on one dataset and adapting it to another similar task.
 4. **Task Similarity:** It's most useful when the tasks are closely related. The model's learned features from the source task are beneficial for the target task.
 5. **Use of Open-Source Libraries:** If the pre-trained model is built using libraries like TensorFlow or PyTorch, you can easily load the
- What is Multi-Task Learning?

Multi-Task Learning (MTL) is a machine learning approach where a single model is trained to perform multiple tasks simultaneously. Instead of building separate models for each task, MTL leverages commonalities across tasks to improve learning efficiency and generalization..



🔗 How Does It Work?

- **Shared Layers:** The model has shared hidden layers that learn general features from all tasks.
- **Task-Specific Layers:** On top of the shared layers, separate output layers are added for each task.
- **Joint Training:** The model is trained on multiple tasks together, using a combined loss function that considers all tasks.

Example:

A neural network trained to do both object detection and image segmentation on the same input image. Shared layers learn visual features, while task-specific layers output bounding boxes or pixel-wise masks.

🚀 When to Use Multi-Task Learning

- When tasks are **related** (e.g., sentiment analysis + emotion detection).
- When you have **limited labeled data** for individual tasks but more combined data.
- When you want to **reduce training time** and computational cost by using one model for multiple purposes.
- When you aim for **better generalization** by leveraging signals from multiple learning objectives.

Advantages of Multi-Task Learning

1. **Improves Generalization:** By learning multiple related tasks at the same time, the model captures shared patterns, leading to better generalization across tasks.
2. **Efficient Learning:** Multi-task learning shares representations among tasks, reducing the need for separate models and making training more efficient.
3. **Reduces Overfitting:** Since the model learns from multiple tasks, it is less likely to overfit on a single task with limited data.
4. **Enables Transfer of Knowledge:** Learning one task can help improve performance on another related task by sharing useful information across tasks.

End-to-End Deep Learning (E2E Learning)

End-to-End Deep Learning is an approach where a deep neural network learns the **entire mapping from input to output** in a single model. Unlike traditional systems where different parts of the task (like feature extraction, classification, decision-making) are handled by separate modules, end-to-end models handle everything **automatically** through a unified architecture.

🔍 Detailed Working

Let's break down the process step-by-step:

1. Input Layer:

The model receives **raw data** (e.g., pixel values of an image, audio waveform, or plain text). There is no need for manual pre-processing or hand-crafted feature extraction.

2. Hidden Layers:

Multiple deep layers (CNNs, RNNs, LSTMs, Transformers, etc.) extract patterns, learn features, and abstract information at different levels:

- For images: Convolutional layers learn spatial hierarchies (e.g., edges, shapes, objects).
- For text/audio: Recurrent or attention-based models capture sequential dependencies.

3. Output Layer:

The final layer directly produces the prediction or decision (e.g., object label, translated sentence, action in robotics, etc.).

4. Loss Function and Optimization:

A loss function evaluates the difference between predicted output and actual output. Backpropagation updates weights across **all layers simultaneously**, optimizing the entire model end-to-end.

💡 Examples

1. **Image Classification:**
 - Input: Raw image (e.g., 224x224 pixels).
 - Output: Predicted class (e.g., "cat" or "dog").
2. **Speech Recognition:**
 - Input: Raw audio waveform.
 - Output: Transcribed text.

Introduction to CNN models: LeNet – 5, AlexNet, VGG – 16, Residual Networks

Unit 6:

Pattern Classification & Recognition Systems - Extended Notes

■ 1. Pattern Classification: Recognition of Olympic Games Symbols

Definition: Pattern classification involves identifying patterns and categorizing them into predefined classes using algorithms trained on visual features.

Use Case: Recognizing various symbols such as the Olympic rings, torches, flags, mascots, and sports pictograms.

Detailed Workflow:

1. **Dataset Collection:**
 - Collect a diverse set of Olympic-related symbols.
 - Include different resolutions, angles, and background conditions.
2. **Preprocessing:**
 - Standardize image size and resolution.
 - Convert color images to grayscale or binary for consistent processing.
 - Apply filters (median, Gaussian) to remove background noise.
3. **Feature Extraction:**
 - **Edge Detection:** Detect outlines and boundaries.
 - **Corner Detection:** Identify sharp changes in image direction.
 - **Texture Analysis:** Analyze surface patterns.
 - **Shape Descriptors:** Zernike moments, HOG (Histogram of Oriented Gradients).
4. **Classification Algorithms:**
 - **Support Vector Machines (SVMs):** Effective in high-dimensional spaces.
 - **Random Forests:** Handles overfitting, interpretable.
 - **Convolutional Neural Networks (CNNs):** Extracts spatial hierarchies of features.
5. **Training & Testing:**
 - Use **cross-validation**.
 - Use **confusion matrix**, **ROC curves**, and **F1-score** for evaluation.

Applications:

- Automated event detection in broadcasts.
- Visual search in Olympic databases.
- Educational and archival tools.

■ 2. Recognition of Printed Characters (OCR)

Introduction: OCR (Optical Character Recognition) translates images of printed text into machine-encoded text.

Core Components:

1. **Image Capture:**
 - Camera/scanner captures text image.
 - Prefer high-resolution and flat-angled images.
2. **Preprocessing Techniques:**
 - **Deskewing:** Align tilted text.
 - **Binarization:** Otsu's thresholding.
 - **Morphological Operations:** Remove noise, bridge gaps.
3. **Segmentation Methods:**
 - Line, word, and character segmentation.
 - Use of projection histograms.
4. **Feature Extraction:**
 - **Geometric:** Height, width, stroke density.
 - **Statistical:** Zoning features, pixel distribution.
 - **Transform-based:** DCT, Fourier descriptors.
5. **Classification Techniques:**
 - **Template Matching:** Limited flexibility.
 - **ML Models:** SVM, MLP.
 - **Deep Learning:** CNNs with LSTM for sequence modeling.
6. **Error Handling:**
 - Use **language models** (n-grams).
 - Contextual correction with NLP.

Real-World Use Cases:

- Bank cheque reading.
- Automated data entry.
- Reading utility bills and receipts.

✍️ 3. Neocognitron: Recognition of Handwritten Characters

Background: Invented by Kunihiro Fukushima in 1979, inspired by the visual cortex.

Network Layers:

- **Input Layer:** Raw 2D pixel input.
- **S-Cells:** Detect primitive features.
- **C-Cells:** Achieve translation invariance.
- **Pooling Mechanism:** C-cells pool responses from S-cells.

Training:

- Unsupervised Hebbian-like learning.
- Gradually increases abstraction layer by layer.

Learning Behavior:

- Early layers: Detect strokes, orientations.
- Later layers: Recognize combinations forming characters.

Applications:

- Postal code and bank form recognition.
- Japanese/Chinese handwriting recognition.

Advantages over Feedforward Networks:

- Better position tolerance.
- Early blueprint for modern CNNs.

🔊 4. NET Talk: Text to Speech (TTS)

History: Developed in the 1980s by Sejnowski and Rosenberg.

Architecture:

- Uses **Time-Delay Neural Network (TDNN)**.
- Trains to map character input to phoneme output.

Working:

1. Character input is encoded in a windowed context.
2. TDNN processes this sequence to predict phonemes.
3. Synthesizer module generates corresponding speech.

Training Data: Large datasets of text aligned with phonetic transcriptions.

Modern Extensions:

- Inspired development of Tacotron, WaveNet, etc.

Applications:

- Screen readers.
- Digital voice assistants (e.g., Siri, Alexa).
- Voiceovers and audiobooks.

🖼️ 5. Recognition of Consonant-Vowel (CV) Segments

Definition: Detecting and distinguishing between consonant and vowel phonemes in speech.

Fundamentals:

- **Vowels:** Harmonic, longer, steady formant patterns.
- **Consonants:** Transient, noise-like, fast-changing.

HMM-Based CV Recognition:

- Each CV pair modeled as a sequence of states.
- Gaussian Mixture Models (GMMs) used to represent acoustic features.

Neural Methods:

- CNN + RNN hybrids for sequential data.
- End-to-end models learn features and classification jointly.

Challenges:

- Speaker variability, accents, background noise.
- Temporal alignment of speech frames.

Evaluation:

- Phone error rate (PER), F1-score.

Applications:

- Language tutoring systems.
- Automated subtitle generation.
- Assistive speech technologies.

🖼️ 6. Texture Classification & Segmentation

Texture Classification:

- **Goal:** Assign a label to an image or image patch based on texture.
- **Common Techniques:**
 - Gabor filters.
 - LBP (Local Binary Patterns).
 - Haralick features.

Texture Segmentation:

- Divide image based on texture similarities.
- **Methods:**
 - Filter bank responses + clustering.
 - Graph-cut optimization.
 - Region merging/splitting.

Advanced Models:

- Deep learning approaches (U-Net, FCN) for end-to-end segmentation.
- Multi-scale CNNs to handle varying texture resolutions.

Use Cases:

- Skin disease detection.
- Satellite imagery classification.
- Automated quality control in manufacturing.